



---

# FIRENET - MODEL COMPRESSION

310552038 蔡瀚興

310554040 葉昭宏

309513017 黃梓熏

葉昭宏 310554040

# What is the problem you want to explore?

在原本的論文中，提出FireNet model，而藉由另一篇論文

FireNet: A Specialized Lightweight Fire & Smoke Detection Model for Real-Time IoT Applications 當中，可以知道FireNet能夠被應用在edge devices上，為了使model的size降低且不太降低accuracy的前提下

採用model quantization+ model pruning修改現有的model



# What is your methodology?

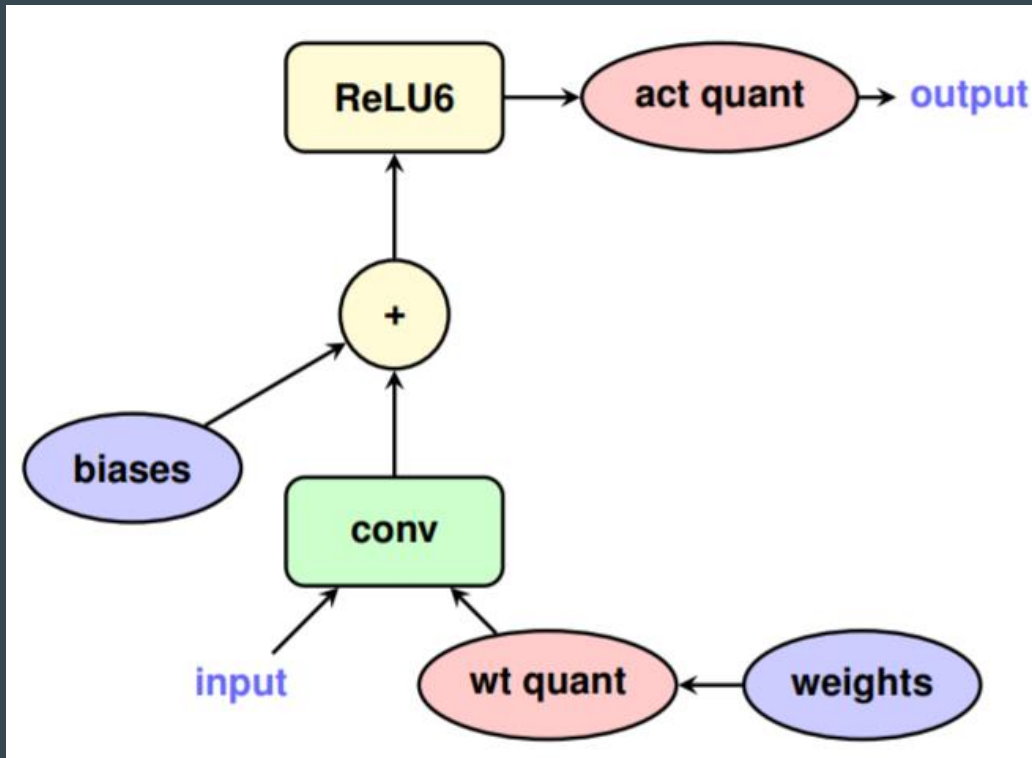
Quantization aware training

use post-training quantization and

quantization-aware training to optimize model

# Introduction to Quantization aware training

In comparison with post-training, Quantization aware training can quantized specific layer



# Goal

1. Make FireNet model smaller than original
1. Make modified model's accuracy same or lower a little bit

# What have you done?

1. test post-training quantization and quantization-aware result
2. test pruned + quantized model and find optimized combination
3. use keras tensorflow\_model\_optimization to control quantized layer

**Result of my work**





# Baseline model detail

Precision: 0.89189

Recall: 0.99

Accuracy: 0.935

Size of gzipped baseline Keras model: 7061894.00 bytes



# Post-training part

Average FPS : 45.31368028167889(so low)

Accuracy : 0.935

Baseline model

Size: 7061894

Size of gzipped post-training TFlite model: 529885.00 bytes

Even if this training model size is the smallest model among many example, for the reason of speed, we didn't take post-training in consideration

# Default quantization-aware training part

Average FPS : 243.55

Accuracy : 0.5

Baseline model

Size: 7061894

Size of gzipped post-training TFlite model: 761781.00 bytes

# Dense layer + 8 bit quantized

Average FPS : 296.10078034062116(moderate speed)

Accuracy : 0.935 (we chose for result)

Size of gzipped 8bitDense TFlite model: 815270.00 bytes

Baseline model

Size: 7061894

# Conv2D layer + 8 bit quantized

Average FPS : 436.98

Accuracy : 0.5

Size of gzipped 8bitConv2D TFlite model: 2352752.00 bytes

Baseline model

Size: 7061894

In this model, we found that the speed is the fastest, but the accuracy is a disaster

# Dense layer + 8 bit quantized + pruned(best pruning)

Average FPS : 291.44

Accuracy : 0.91

Size of gzipped 8bitConv2D TFlite model: 982336.00 bytes

Baseline model

Size: 7061894

# Conclusion in quantization

Best combination in quantization:

Dense layer only + 8bit quantization (moderate speed + moderate accuracy)

(we will show result in raspberry pi part)

There is a trade-off between pruning and quantization

310552038 蔡瀚興



# What is the problem you want to explore?

原本論文提出的Firenet model，效率及準確率都還可以改進。

在Google Colab中跑論文給的model，結果如下：

FPS  $\approx$  12 (per sec)

Accuracy  $\approx$  0.935

Model weight size  $\approx$  7061894.00 (bytes)

因此我們希望能夠找出方法，增加FPS及縮小Model weight size，甚至提高Accuracy。

# What is your methodology?

## Magnitude-based weight pruning

- After each training epochs, the link with the smallest weight is removed.
- Thus the saliency of a link is just the absolute size of its weight.
- Though this method is simple, it rarely yields worse results than the more complicate algorithms.

# What have you done?

1. Find dataset and build the Colab training environment
2. Write the python script for testing result on Colab and Raspberry Pi.
3. Find & try different model pruning strategy to optimize the result.

**Result of my work**

# BASLINE MODEL

## Hyperparameter

Batch = 32

Epochs = 100

Loss = sparse\_categorical\_crossentropy

optimizer = adam

## Dataset

Fire : 1124

NoFire : 1301



Layer (type)	Output Shape
conv2d (Conv2D)	(None, 62, 62, 16)
average_pooling2d (AveragePooling2D)	(None, 31, 31, 16)
dropout (Dropout)	(None, 31, 31, 16)
conv2d_1 (Conv2D)	(None, 29, 29, 32)
average_pooling2d_1 (AveragePooling2D)	(None, 14, 14, 32)
dropout_1 (Dropout)	(None, 14, 14, 32)
conv2d_2 (Conv2D)	(None, 12, 12, 64)
average_pooling2d_2 (AveragePooling2D)	(None, 6, 6, 64)
dropout_2 (Dropout)	(None, 6, 6, 64)
flatten (Flatten)	(None, 2304)
dense (Dense)	(None, 256)
dropout_3 (Dropout)	(None, 256)
dense_1 (Dense)	(None, 128)
dense_2 (Dense)	(None, 2)

# BASELINE MODEL

Precision: 0.89189


Recall: 0.99

Accuracy: 0.935



Size of gzipped baseline Keras model:  
7061894.00 bytes

# PRUNING ONLY DENSE LAYER



Precision: 0.93333

Recall: 0.98

Accuracy: 0.955

Baseline model


Accuracy: 0.935

Size of gzipped baseline Keras model : 7061894.00 bytes

Size of gzipped pruned Keras model : 1536520.00 bytes

Layer (type)	Output Shape
conv2d (Conv2D)	(None, 62, 62, 16)
average_pooling2d (AveragePooling2D)	(None, 31, 31, 16)
dropout (Dropout)	(None, 31, 31, 16)
conv2d_1 (Conv2D)	(None, 29, 29, 32)
average_pooling2d_1 (AveragePooling2D)	(None, 14, 14, 32)
dropout_1 (Dropout)	(None, 14, 14, 32)
conv2d_2 (Conv2D)	(None, 12, 12, 64)
average_pooling2d_2 (AveragePooling2D)	(None, 6, 6, 64)
dropout_2 (Dropout)	(None, 6, 6, 64)
flatten (Flatten)	(None, 2304)
prune_low_magnitude_dense (PruneLowMagnitude)	(None, 256)
dropout_3 (Dropout)	(None, 256)
prune_low_magnitude_dense_1 (PruneLowMagnitude)	(None, 128)
prune_low_magnitude_dense_2 (PruneLowMagnitude)	(None, 2)

# PRUNING ONLY CONV LAYER



Precision: 0.93458

Recall: 1.0

Accuracy: 0.965

Baseline model

Accuracy: 0.935

Layer (type)	Output Shape
prune_low_magnitude_conv2d (PruneLowMagnitude)	(None, 62, 62, 16)
average_pooling2d (AveragePooling2D)	(None, 31, 31, 16)
dropout (Dropout)	(None, 31, 31, 16)
prune_low_magnitude_conv2d_1 (PruneLowMagnitude)	(None, 29, 29, 32)
average_pooling2d_1 (AveragePooling2D)	(None, 14, 14, 32)
dropout_1 (Dropout)	(None, 14, 14, 32)
prune_low_magnitude_conv2d_2 (PruneLowMagnitude)	(None, 12, 12, 64)
average_pooling2d_2 (AveragePooling2D)	(None, 6, 6, 64)
dropout_2 (Dropout)	(None, 6, 6, 64)
flatten (Flatten)	(None, 2304)
dense (Dense)	(None, 256)
dropout_3 (Dropout)	(None, 256)
dense_1 (Dense)	(None, 128)
dense_2 (Dense)	(None, 2)

Size of gzipped baseline Keras model : 7061894.00 bytes

Size of gzipped pruned Keras model : 2377961.00 bytes



# PRUNING ALL LAYER

## Three pruning strategy :

- PolynomialDecay from 100% ~ 70%
- PolynomialDecay from 100% ~ 20%
- ConstantSparsity with 50%

Layer (type)	Output Shape
prune_low_magnitude_conv2d (PruneLowMagnitude)	(None, 62, 62, 16)
prune_low_magnitude_average_pooling2d (PruneLowMagnitude)	(None, 31, 31, 16)
prune_low_magnitude_dropout (PruneLowMagnitude)	(None, 31, 31, 16)
prune_low_magnitude_conv2d_1 (PruneLowMagnitude)	(None, 29, 29, 32)
prune_low_magnitude_average_pooling2d_1 (PruneLowMagnitude)	(None, 14, 14, 32)
prune_low_magnitude_dropout_1 (PruneLowMagnitude)	(None, 14, 14, 32)
prune_low_magnitude_conv2d_2 (PruneLowMagnitude)	(None, 12, 12, 64)
prune_low_magnitude_average_pooling2d_2 (PruneLowMagnitude)	(None, 6, 6, 64)
prune_low_magnitude_dropout_2 (PruneLowMagnitude)	(None, 6, 6, 64)
prune_low_magnitude_flatten (PruneLowMagnitude)	(None, 2304)
prune_low_magnitude_dense (PruneLowMagnitude)	(None, 256)
prune_low_magnitude_dropout_3 (PruneLowMagnitude)	(None, 256)
prune_low_magnitude_dense_1 (PruneLowMagnitude)	(None, 128)
prune_low_magnitude_dense_2 (PruneLowMagnitude)	(None, 2)

# PolynomialDecay from 100% ~ 70%


Precision: 0.95652

Recall: 0.88

Accuracy: 0.92

Baseline model

Accuracy: 0.935

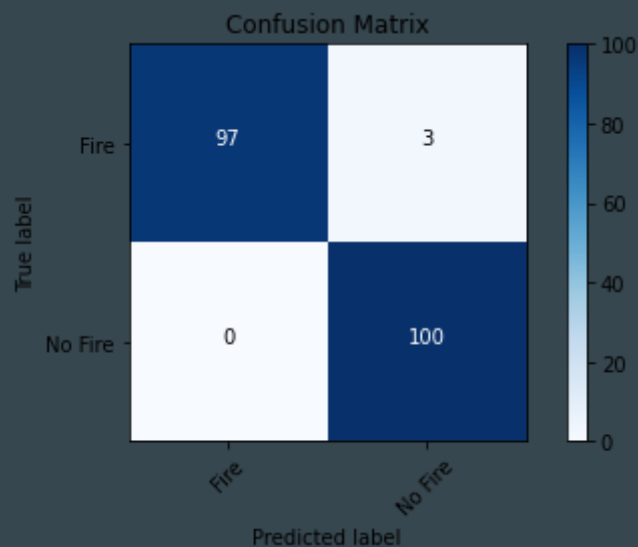


- Size of gzipped baseline Keras model : 7061894.00 bytes
- Size of gzipped pruned Keras model : 774873.00 bytes

# PolynomialDecay from 100% ~ 20%

Precision: 1.0  
Recall: 0.97  
Accuracy: 0.985

Baseline model  
Accuracy: 0.935



- Size of gzipped baseline Keras model : 7061894.00 bytes
- Size of gzipped pruned Keras model : 772341.00 bytes

# ConstantSparsity with 50%


Precision: 0.85714

Recall: 0.96

Accuracy: 0.9

Baseline model

Accuracy: 0.935



- Size of gzipped baseline Keras model : 7061894.00 bytes
- Size of gzipped pruned Keras model : 1513879.00 bytes

309513017 黃梓熏

# What is the problem you want to explore?

嘗試將Model Pruning與Quantization結合，並同時套用到Firenet model上，來觀察是否能讓model在保持與Baseline model(只單用Model Pruning的model)差不多的準確度的情況下，model的size有大幅降低的效果。

# What is your methodology?

Pruning preserving quantization aware training (PQAT)

→ 是tensorflow底下提供的一種優化方法，能讓做過pruning的model在維持原本pruning的效果下再去做quantization。

# What have you done?

1. Combine pruning and quantization into Firenet model by using PQAT
2. Setup the execution environment for Raspberry Pi 4b
3. Test the performance of tensorflow lite model on Raspberry Pi 4b



**Result of my work**

# PQAT

**Pruning method : PolynomialDecay from 100% ~ 20%**

**Quantization method : 8 bits prune preserve quantization**

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 62, 62, 16)	448
average_pooling2d (AveragePooling2D)	(None, 31, 31, 16)	0
dropout (Dropout)	(None, 31, 31, 16)	0
conv2d_1 (Conv2D)	(None, 29, 29, 32)	4640
average_pooling2d_1 (AveragePooling2D)	(None, 14, 14, 32)	0
dropout_1 (Dropout)	(None, 14, 14, 32)	0
conv2d_2 (Conv2D)	(None, 12, 12, 64)	18496
average_pooling2d_2 (AveragePooling2D)	(None, 6, 6, 64)	0
dropout_2 (Dropout)	(None, 6, 6, 64)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 256)	590080
dropout_3 (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 2)	258

=====  
Total params: 646,818  
Trainable params: 646,818  
Non-trainable params: 0

# Model Size

- Size of gzipped origin Keras model : 7061894.00 bytes
- Size of gzipped pruned Keras model (baseline) : 772341.00 bytes
- Size of gzipped pruned TFlite model(baseline) : 754639.00 bytes
- Size of gzipped pruned TFlite model(with PQAT) : 225301.00 bytes

# Accuracy

	Dataset 1	Dataset 2	Dataset 3
Baseline model (Keras)	0.92	0.985	0.965
Baseline model (TFlite)	0.92	0.985	0.965
Model with PQAT (TFlite)	0.965	0.965	0.955

**Result of our work**

# Baseline model

- We can't run baseline model on Raspberry Pi is too lag
- The next page is the result of real time video run on Colab

# Baseline model

```
CPU usage : 100.0 %  
RAM memory: 7.97 %  
Time taken = 0.10644769668579102  
FPS: 9.394284997558648  
Fire Probability: 100.0  
(1, 64, 64, 3)
```

Fire Probability: 100.0



# Best pruning strategy – PolynomialDecay from 100% ~ 20%


```
CPU usage : 100.0 %  
RAM memory: 12.62 %  
Time taken = 0.003481626510620117  
FPS: 287.22207765527634  
Fire Probability: 99.99957084655762  
(1, 64, 64, 3)
```

Fire Probability: 99.99957084655762





# Best pruning stratgy – PolynomialDecay from 100% ~ 20%



<http://drive.google.com/file/d/12SXF162-zVikSd3rwXMuh2kLwLngsqbS/view>


# Best quantization stratgy – int 8 bit on dense layer only

```
CPU usage : 81.2 %  
RAM memory: 9.31 %  
Time taken = 0.00715327262878418  
FPS: 139.79615371796154  
Fire Probability: 100.0  
(1, 64, 64, 3)
```

Fire Probability: 100.0



# Best quantization stratgy – int 8 bit on dense layer only



[http://drive.google.com/file/d/143fAWxbA8UcbP97n57CdQ2Rg7\\_BnZsJ7/view](http://drive.google.com/file/d/143fAWxbA8UcbP97n57CdQ2Rg7_BnZsJ7/view)

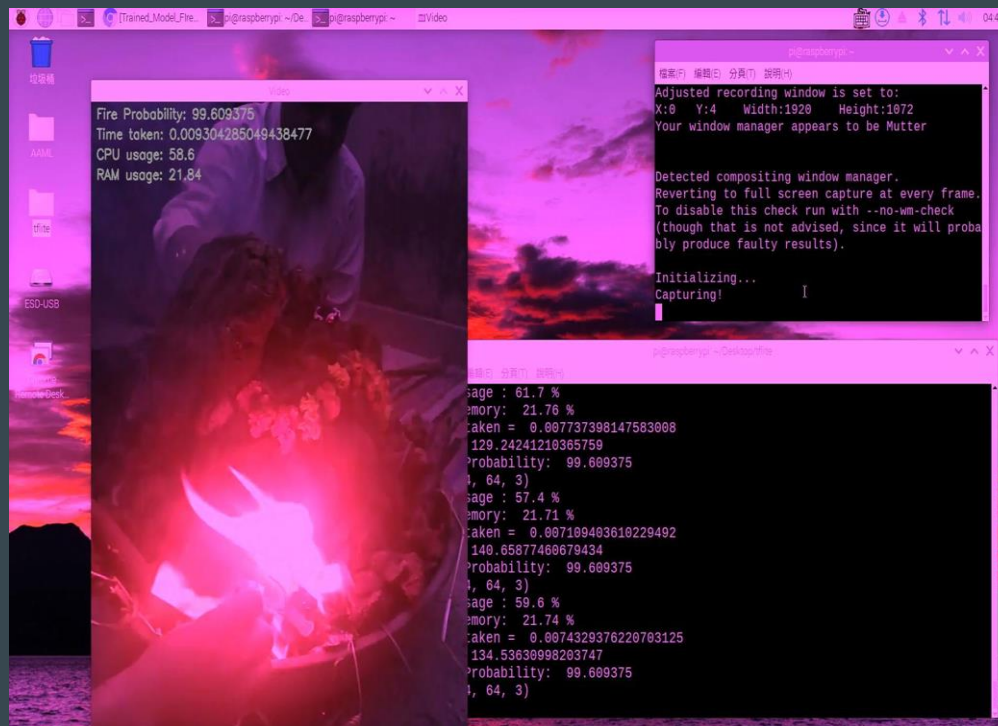
# Do best pruning strategy first then do best quantization strategy

```
CPU usage : 57.4 %  
RAM memory: 8.25 %  
Time taken = 0.0030732154846191406  
FPS: 325.39208688906126  
Fire Probability: 99.609375  
(1, 64, 64, 3)
```

Fire Probability: 99.609375



# Do best pruning then best quantization



<http://drive.google.com/file/d/1r18K1xDv7iAJEVCTrydJhLCgP2JJU5mM/view>


# PQAT : PolynomialDecay from 100% ~ 20% + 8 bits prune preserve quantization

```
CPU usage : 100.0 %  
RAM memory: 12.22 %  
Time taken = 0.049668312072753906  
FPS: 20.133561183541023  
Fire Probability: 99.609375  
(1, 64, 64, 3)
```

Fire Probability: 99.609375



**PQAT** : PolynomialDecay from 100% ~ 20% + **8 bits prune preserve quantization**



<http://drive.google.com/file/d/14qsNR3IXSFZVGxHdBL5GVJ96HlanJEFO/view>